COLLIN O'NEIL

Lying, Trust, and Gratitude

I. INTRODUCTION

Deception can be achieved by a wide variety of methods. In the most recent primary election season, observers identified twelve distinct deceptive techniques deployed in political advertisements, ranging from out-of-date evidence and selective editing to photoshopping and deceptive dramatization.[1] But it is deception via the hoary technique of lying that seems to receive the bulk of the attention and condemnation in both our public culture and our private lives. A recent Pulitzer Prize–winning book that relates the lies of public figures like Martha Stewart and Bernard Madoff is subtitled "How False Statements Are Undermining America."[2] Fact-checking organizations like Politifact have taken on an increasingly prominent role in policing political discourse; yet the facts they check are limited to those the politicians state, not what the politicians may intend the public to believe via other methods.

The preoccupation with lying may derive in part from the fact that it is easier to attribute a deceptive intention to someone deploying a false

An earlier version of this article was presented at the Conference on Deception and Trust at the University of Sheffield, and I thank the organizer, David Owens, and the other participants for thought-provoking discussion. I am especially indebted to Seana Shiffrin, who encouraged my interest in this topic and has offered many valuable suggestions along the way. Many thanks are also due to Richard Arneson, Barbara Herman, Pamela Hieronymi, Arnon Keren, Chris Kochevar, Chiara Lepora, Franklin Miller, Joseph Millum, Marsha O'Neil, David Owens, William Ruddick, Benjamin Sachs, and Alan Wertheimer for very useful comments on drafts at various stages. Comments from two anonymous referees on later drafts prompted significant revisions and clarifications, and I am grateful for their advice and criticisms.

1. <http://www.flackcheck.org/patterns-of-deception/> Accessed October 15, 2012.

2. James B. Stewart, *Tangled Webs: How False Statements Are Undermining America* (New York: Penguin, 2011).

statement than some other mode of deception. But it also reflects the common moral judgment that lying is in a way worse than other forms of deception: the lie is a "special affront." Philosophers, by and large, have been sympathetic to the idea that the lie is special. But it is puzzling that the contrast they have seemed most concerned to establish has been the contrast between deceiving by saying something false and deceiving by saying something true that implies something false, that is, through false implicature.[3] There is another method of deception, "covert deceit," that has a greater prima facie claim to special status than both the lie and false implicature.

Suppose that Bob spills coffee on Sue's laptop and wants to avoid having to pay to replace it. One thing he might do is lie to her when she returns home, saying that the cat knocked the coffee over when he was in the bathroom. Or he might remind her (truthfully) that he swore off coffee two weeks ago, implying that he was not responsible, even though he relapsed that morning. But a rather different thing he might do is dip the cat's paws in cold coffee so that when Sue returns home she will see those stained paws and conclude that the cat did it. Whether Bob lies to Sue or misleads her by saying something true, he is at least open with her about his intention to get her to believe that he did not do it. But he keeps this intention hidden from Sue when he frames the cat, which would appear to make this method rather more manipulative. Staining the cat's paws may be a contrived example, but covert methods of deceit are not uncommon: wearing a fake Rolex to appear wealthy, forging documents, planting a weapon on a body, displaying a counterfeit Medal of Honor, crying crocodile tears, holding off sending a work-related e-mail until late at night to give the false impression that one has been working long hours, and so on.

Covert methods present a challenge for the idea that the lie is a special affront. How might the greater openness of the lie, which makes it less devious and thus in a clear way better than covert methods, somehow make it in another way worse? The current philosophical emphasis on the contrast between lying and false implicatures, and the relative

3. See, e.g., Jonathan Adler, "Lying, Deceiving, or Falsely Implicating," *Journal of Philosophy* 94 (1997): 435–52; and Bernard Williams, *Truth and Truthfulness* (Princetonn N.J.: Princeton University Press, 2002), pp. 84–110.

neglect of the contrast between lying and covert methods, may be reflective of the following line of thought.

The lie is open, which makes it less devious, but it is open in this way because it invites a kind of trust from the victim, and this invitation to trust, in turn, puts the deceiver into a morally problematic relation to the victim's trust. Covert methods obviously do not invite trust, and therefore never enter into this problematic relation to the victim's trust. Thus, while the invitation to trust, and the openness about the intent to produce a belief that follows from it, may make the lie less manipulative, the lie is still manipulative. Moreover, it commits a different *kind* of wrong than covert deceit. Since it is not obvious, however, that false implicatures never invite trust, the account that separates the lie from covert deceit may not separate lies from false implicatures. So, the contrast between lies and false implicatures still demands philosophical attention.

There is something correct in this line of reasoning, and I will ultimately try to vindicate it, though only partially. But it is incomplete as it stands: so much depends on the details of the account of the morally problematic relation to trust. Nearly everyone would agree that trust can be morally significant—we have familiar names for the distinctive wrongs that involve trust, like "betrayals" and "abuses" of trust. But I do not think that we have anything like an adequate understanding of the significance of trust. A major aim of this article is to provide a plausible account of its significance. Although "betrayals" and "abuses" are often meant interchangeably, I think they correspond to two different morally problematic relations to trust. I will try to explain the nature of these wrongs by appeal to the idea that we should be *grateful* for trust. Betrayals and abuses of trust violate distinct demands of gratitude for trust.

The accounts I provide of betrayals and abuses of trust can be usefully applied to the question of whether the lie is indeed special in virtue of its connection to trust. The special wrongness of the lie, as compared with covert deceit, cannot be made out as a betrayal of trust. But it can be explained as an abuse of trust. However, the lie (and its cousin, the false implicature) is not unique in its capacity to abuse trust. There is another familiar method of deceit, "showing," that, unlike covert methods, is open about the intention to produce a belief but, unlike the lie, does not invite trust. Deceit by showing is on a par with the lie. This shows that the

morally fundamental category of deceit, as far as trust is concerned, is not the lie but rather deceit by means of communication.

II. LYING AND THE INVITATION TO TRUST

Lies invite trust, in the sense that they invite the audience to take the speaker's word for the truth of their content.[4] This simple characterization of the *invitation to trust* will serve for now. The presence of an invitation to trust in this sense may not distinguish the lie from all false implicatures. But covert methods of deception, such as staining the cat's paws, never invite trust, and this makes the invitation to trust an appealing hook on which to hang an account of the specialness of the lie.

There are two different ways in which the significance of the invitation to trust has been understood: as setting up a special obligation of some sort that the lie breaks, or as marking an intention to use a certain mechanism for the deceit—the mechanism of trust—that somehow makes the act of deceit morally worse. I will call these different approaches the "Special Obligation Approach" and the "Use Approach." The most prominent version of the Special Obligation Approach claims that the speaker's invitation to take his word for the content he communicates is really an invitation to accept an implicit promise that the content is true. Thus the liar, in asserting something false, breaks a promise to the hearer.[5] Supposing that lies do make such a promise would explain why Bob's lie compares unfavorably with staining the cat's paws: although both methods of deception wrong Sue, the lie also makes and breaks a promise to her.

But it is doubtful that an invitation to take a speaker's word for something presupposes such a promise. If this were the case, then whenever a speaker tells someone a falsehood, the speaker's behavior is morally on a par with that of someone who breaks a promise, even when the

4. I do assume that lies invite trust, but I do not mean to imply that lies must be intended to deceive. See Thomas Carson, *Lying and Deception: Theory and Practice* (Oxford: Oxford University Press, 2010), pp. 20–23, for the view that lying need not be deceptive. Inviting trust requires only that one manifest an intention that the audience form a trusting belief, not that one have this intention. However, my topic is methods of *deceit*, so when I talk about lies I am always referring to deceptive lies.

5. See W. D. Ross, *The Right and the Good* (Oxford: Oxford University Press, 1930), p. 21; Thomas Carson, "The Definition of Lying," *Nous* 40 (2006): 284–306; and Charles Fried, *Right and Wrong* (Cambridge, Mass.: Harvard University Press, 1978), p. 67.

assertion he makes is sincere and adequately supported. Yet when someone makes an assertion of this kind, there is normally no recrimination if it turns out to be false. Some kind of apology may be in order, but it would be understood as an admission of a mistake, not as an admission of a wrong.[6]

The Use Approach assigns a rather different significance to the invitation to trust. The invitation to trust is thought to make a difference, not because it triggers some special obligation that the lie violates, but because it marks an attempt to *use* the victim's trust to get them to believe something. The account I will eventually defend takes this form, but first I will relate another version that is unsuccessful.

According to this version, when a speaker invites trust, he is intending to induce a false belief via a noninferential process, but when he employs a covert method, like staining the cat's paws, the victim arrives at the false belief only via an inference from the evidence he has planted. This in turn is thought to imply that the victim of the lie is not responsible for forming the false belief, whereas the victim of the covert method is considered to be responsible.[7]

There are simply too many gaps in this argument for it to be convincing. Lies can induce false beliefs noninferentially, although this is at most true of only some lies. But why would it follow from the fact that a victim arrives at a belief without inferring it that she is not responsible for believing falsely? And even if this did follow, it is unclear why this would make lying worse than planting misleading evidence. One might just as plausibly take the fact that planting evidence makes the victim unwittingly complicit in their own deception as a reason to think that such a method is worse than the lie, not better.[8]

6. There is a version of the promissory account that would avoid this implication, namely, that the lie makes a promise that the content of the assertion is believed, not true. This may be the view in Roderick M. Chisholm and Thomas D. Feehan, "The Intent to Deceive," *Journal of Philosophy* 74 (1977): 153. However, both this version and the version according to which the lie promises that the content is true have trouble explaining why a promise not to lie on some topic does not always seem redundant in the way it should if assertions implicitly made such a promise.

7. See, e.g., Stuart P. Green, "Lying, Misleading, and Falsely Denying," *Hastings Law Journal* 53 (1997): 154; and see also Jonathan Adler, "Lying, Deceiving, or Falsely Implicating," p. 444.

8. Alan Strudler may also be offering a version of the Use Approach in Strudler, "Deception Unraveled," *Journal of Philosophy* 102 (2005): 458–73. He argues that deception that

The foregoing accounts of the special significance of the invitation to trust are in my opinion unsuccessful. But I do think the lie's invitation to trust has the potential to aggravate the wrong of deception. In addition to whatever is wrong with deception itself, inviting trust as the lie does can put the deceiver into a certain morally problematic relation to the victim's trust, a relation that covert methods of deception never bear to trust. To make my case I will need to distinguish two different problematic relations that a deceiver, or a wrongdoer more generally, may bear to a victim's trust and to explain precisely why each of these relations is problematic. One relation consists in *failing to perform as one is trusted to perform*. The associated wrong is a "betrayal of trust." The other relation consists in *misusing someone's trust*, and the associated wrong is an "abuse of trust."

Robert Solomon appeals to the first relation in explaining what is special about the lie. He claims, "A lie breaches the very trust it necessarily presupposes . . . in the very personal and concrete sense that usually goes by the name of 'betrayal.' "[9] This idea bears a similarity to the promissory account described earlier, in that the invitation to trust is thought to trigger a special moral obligation that the lie violates. Although the lie does not make and break a promise, it might still set up and commit a betrayal. But there is an obstacle to evaluating this proposal. At least we know what promise-breaking is; the same cannot be said of a betrayal of trust. In the following section, I address this gap by

---

makes use of a certain kind of trust is special because it is more manipulative than deception that does not. The reason, he says, is that when you trust a speaker, "you have transferred the effective locus of your decision-making on your beliefs . . . to me." As will emerge later, there is an important point of agreement between his account and my own: we both take the kind of trust that is relevant to the question of whether some methods of deceit are worse than others to be the kind involving an expectation of good will. But his reasoning about the significance of using that trust appears to overlook the problem of covert deceit mentioned at the outset of this article. When I invite your trust and you trust me, your beliefs are dependent on my will. However, your beliefs are dependent on my will when I covertly deceive you as well, but you just do not know it. The covert deception seems to be more manipulative, not less. See also Paul Faulkner, "What Is Wrong with Lying?" *Philosophy and Phenomenological Research* 75 (2007): 535–57, for a different kind of attempt to explain what makes the lie manipulative in terms of its connection to trust.

9. Robert Solomon, "What a Tangled Web: Deception and Self-Deception in Philosophy," in *Lying and Deception in Everyday Life*, ed. Michael Lewis and Carolyn Saarni (New York: Guilford Press, 1993), pp. 40–41.

providing an account of betrayals of trust. Lies can commit betrayals of trust, but we will see in the end why this fact cannot satisfactorily explain the specialness of the lie.

III. BETRAYALS OF TRUST

What is a betrayal of trust? Solomon does not say. There is an obvious answer to the question: to betray trust is to fail to do what one is trusted to do. But this answer is unsatisfying. It is far from clear that everything that goes by the name of "trust" makes possible a genuine betrayal. And even if we knew what kind of trust makes betrayals possible, and restricted the obvious answer accordingly, there remains the question of how trust makes this moral difference—what kind of reason it generates. We are familiar with the charge "How could you do that to me? I *trusted* you," which cites trust as somehow adding to the wrong. But we lack an account of what kind of trust adds a reason not to wrong the truster, and what kind of reason this is.

I will provide an account of trust that connects trust to gratitude, thus making sense of another familiar expression of betrayal: "I trusted you, and this is how you *repay* me?" My strategy in developing an account of betrayals of trust will be to divide the problem into two parts. I will first provide a description of a kind of trust that, if any kind of trust does, makes possible a genuine betrayal, namely, the kind of trust that makes us ready to *feel betrayed* upon the discovery that our trust was misplaced.[10] Next, I will examine the components of the kind of trust that makes us ready to feel betrayed to see what kind of moral reason they might generate that a betrayal of trust disregards.

10. Although the literature on trust uses susceptibility to reactive attitudes to distinguish trust from mere reliance, it has not distinguished between the feeling of resentment and the feeling of betrayal. But these are very different moral emotions, and attention to these differences will ultimately prove to be critical in developing an account of betrayals of trust. See Annette Baier, "Trust and Antitrust," *Ethics* 96 (1986): 231–60; Karen Jones, "Trust as an Affective Attitude," *Ethics* 107 (1996): 4–25; Richard Holton, "Deciding to Trust, Coming to Believe," *Australasian Journal of Philosophy* 72 (1994): 63–76. It turns out that will-based accounts of trust, like Baier's and Jones's (and unlike Holton's), are most plausible as accounts of a kind of trust that rationalizes a readiness to feel betrayed, not merely resentful, although this is not how Baier and Jones conceive of their own accounts.

A. The Feeling of Betrayal

I assume that trust requires an expectation of performance. What needs to be added to this expectation to explain a readiness to feel resentment? What then needs to be added to that to explain a readiness to feel betrayed? To begin with a familiar illustration: Kant's neighbors may have expected him to take his daily walk at precisely noon and even used his habits as their clock, but they did not trust him to continue because they were not ready to resent him or reproach him were he to change his habits or oversleep.[11] Had the neighbors trusted Kant to walk at noon each day, they not only would have expected him to do it, but also would have regarded him as obligated to do it.[12] This is what was missing from the way Kant's neighbors regarded his daily walk and explains why they were not prepared to resent him for missing it.

But there is another important kind of trust that not only disposes the truster to feel resentful when the trustee disappoints them, but also prompts the distinct feeling of *betrayal*. Although this kind of trust also involves regarding the trustee as under an obligation, this feature cannot explain why we are ready to feel betrayed, and not merely resentful, by the trustee's failure. These are different ways of experiencing a wrong.

One characteristic of the feeling of betrayal is that it involves an element of surprise. The feeling of resentment or anger at a wrong, by contrast, is compatible with having expected it to occur. But what is the element of surprise in the feeling of betrayal about? It is about the wrong, but is it just the fact that the wrong occurred, or is it something else? I suggest that the surprise in the feeling of betrayal is focused, not primarily on the fact that the wrong occurred, but on something that the occurrence of the wrong reveals. When we feel betrayed by a wrong, we are *hurt* by it and disposed to *rethink* our relationship with the one who wronged us. Being wronged unexpectedly is not always hurtful in the way that characterizes the feeling of betrayal. For example, I may have expected my roommate not to steal from me, but only because I thought

11. Baier, "Trust and Antitrust," p. 235.

12. This may be a nonmoral obligation, in the sense of the highest standard that both the truster and the trustee hold themselves accountable to in the context, for example, an obligation generated by an agreement between two criminals not to talk to the police. I see it as a potential advantage of the gratitude account of betrayals I go on to offer that it could explain why breaking such an agreement, though not impermissible, would be a genuine betrayal of trust.

that she was unaware of the value of my property. When she does indeed steal my property, yes, I will be surprised about it, but I will not be making the painful discovery that I am less important to her than I had thought. Nor will I take myself to have reason to reevaluate my relationship with her. At most, her wrong reminds us of something I already knew about the way she regarded me.

So in the case of a betrayal of trust, we are surprised not just by the fact that the trustee wronged us, but by the fact that the trustee *would* wrong us. This is the key insight of will-based theories of trust, namely, that the readiness to feel betrayed rests on an expectation that the trustee has good will toward us.[13] When we expect not to be wronged by someone only because we think wronging us would run contrary to the person's own interests, this indicates that we *already* do not believe that the trustee has good will toward us or that her good will toward us is sufficiently strong to keep her from wronging us. Thus, if she does wrong us, this is not felt as keenly; it merely confirms what we already believed about our importance to her.

IV. THE NATURE OF A BETRAYAL OF TRUST

The feeling of betrayal is distinguished from the feeling of resentment by the fact that a feeling of betrayal hurts and prompts us to reassess our relationship. But these feelings differ from each other, not merely as ways of experiencing a wrong, but as *blame*. As mentioned earlier, when we are wronged, we may complain, "How could you do that to me?"; but when we also feel betrayed, the complaint naturally takes the form, "How could you do that to me? I trusted you," citing our trust as somehow amplifying the wrong. The sense of surprise and hurt that characterize the feeling of betrayal are not its distinctive moral content

13. By "good will," I do not mean that the trustee's will is benevolent, but merely that it is responsive to the trustee's obligations. We can trust those who do not care how things go for us—we can even trust enemies who hate us—so long as we expect them to be responsive to our rights. We will not be *as* hurt by a wrong when we expected good will only in this narrow sense, but the discovery that an enemy does not even respect one's rights is still hurtful. In cases of paternalism, where the trustee violates the rights of the truster from benevolence, the trustee displays good will in one familiar sense but not in the sense I am discussing. Paternalistic wrongs can still be experienced as betrayals, but the fact that there is good will in the sense of benevolence helps explain some of the ambivalence in the victim's reaction. Not all betrayals are equal.

—these are feelings about the wrong and what it reveals, but they do not add anything to the complaint. The next step is to identify this distinctive moral content.

How should that be done? The kind of trust that makes sense of the feeling of betrayal must also, unless the distinctive moral content of the feeling is systematically misapplied, make the moral difference that satisfies that distinctive moral content. What is left to determine, then, is how the features of this kind of trust can, individually or in combination, give a trustee a special moral reason not to wrong the truster. To identify this reason is to explain what a betrayal of trust is.

In attempting to identify this reason, I do not mean to ignore the possibility that, once we know what kind of trust makes possible a betrayal of trust, there is nothing left to be said—that is, that the reasons this kind of trust generates may be nonreducible to reasons of other kinds. This is, after all, a position that many take with respect to the kind of reasons that special relationships like friendships generate.[14] If a betrayal of loyalty consists in disregarding a nonreducible reason generated by a special relationship, perhaps a betrayal of trust consists in disregarding a nonreducible reason generated by trust. This cannot of course be the same reason. We can trust someone, such as a stranger or even an enemy, without standing in a special relationship to them. And we can stand in a special relationship to someone without trusting the person to honor all of their commitments. A friend's failure to keep even those commitments we do not trust him to keep does remain a betrayal of loyalty. Still, special relationships and trust might generate distinct nonreducible reasons that underlie the wrongs in betrayals of loyalty and betrayals of trust.

In the case of trust, however, such a nonreductive account should be resisted. There is a particular kind of reductive account that can explain how the fact that a wrong disappoints trust can make it a different kind of wrong. And once the component attitudes of the kind of trust that makes possible betrayals are laid bare, this reductive account becomes intuitively compelling. To review, the components of the kind of trust that make us ready to feel betrayed are: (a) expecting that the trustee will discharge (or will try to discharge) her commitment to us, and (b)

14. See Samuel Scheffler, "Relationships and Responsibilities," *Philosophy & Public Affairs* 26 (1997): 189–209.

expecting that she will do this (or will try) because she has good will toward us. To hold these attitudes toward someone whom we regard as under an obligation to us is to positively affirm their will—to have a good opinion of them. This is a benefit,[15] and a special kind of benefit: to trust someone in this way is to *honor* them.[16]

Sometimes we honor others through action, by being willing to marry them or by giving them an important mission or by voting for them in an election. But we can also honor others without taking any action at all. They can be honored by our attitudes in themselves. For the attitude of trust to bestow an honor, both of its elements must be in place. There is no honor if the truster thinks you will do your duty only because it aligns with your own interests.

Thus, whether a truster is ready to feel betrayed by a trustee appears to turn on whether or not their action or attitude honors the trustee. But what distinctive kind of wrong consists in disregard of an honor? The answer, I propose, is the wrong of ingratitude. We should be grateful to be honored. What happens then in a betrayal of trust is a repudiation of that honor—a failure to live up to or make oneself deserving of the honor implicit in trust—and this repudiation is a form of ingratitude.

V. POTENTIAL OBJECTIONS

Even if it is granted that trust generates a reason to do as one is trusted, it might be argued that this reason cannot be a reason of gratitude: a betrayal of trust is something specific, yet debts of gratitude are indefinite; trust can be owed to the trustee, and therefore cannot generate a debt of gratitude; trust can be unwelcome, and therefore cannot create a

15. Philip Pettit calls this an "attitude-dependent good" and points out that, since a trustee will want to keep possession of our good opinion, manifesting trust in someone can give us a reason to expect her to do what we trust her to do without having any evidence of her trustworthiness. Pettit, "The Cunning of Trust," *Philosophy & Public Affairs* 24 (1995): 202–25.

16. When one is expected to meet standards that are hard to meet, the word "honor" seems fitting, but when one is expected to meet, say, standards of common decency, "honor" may seem too inflated. In calling trust an "honor" regardless of how easy it is to meet the relevant standards, I may be stretching the meaning of the word. But I am primarily interested in marking the *kind* of benefit one is bestowing when one affirms the character or abilities of another. Whether the standards are hard or easy, trust positively affirms the trustee's will, and thus bestows this kind of benefit.

reason of gratitude. I will address these objections, and along the way will fill in some important details of the gratitude account.

A. Objection: Debt of Gratitude

Although we should be grateful for an honor and should thus incur some sort of debt, it is not yet clear what gratitude for trust demands. Ordinarily we enjoy considerable latitude when it comes to discharging debts of gratitude. We should return favors, but when and how we should do so is indefinite. The fact that reasons of gratitude are ordinarily indefinite has been cited against gratitude accounts of political obligation—the obligation to obey the law. Even if we should be grateful for benefits we receive from the state, why would obedience to the law be the uniquely suitable return? Yet I am claiming that a betrayal of trust—a failure to do what one is trusted to do—is ungrateful. For this to be the case, gratitude for trust must demand something very specific, namely, doing what one is trusted to do. If you fail in this, you are ungrateful for the trust, even if you are ready to make some other form of return. There is no latitude in how you express your gratitude.

Indeed, a gratitude account of betrayals of trust might appear to have even greater difficulty handling this objection than gratitude accounts of political obligations. A.D.M. Walker has argued, in defense of his own gratitude account of political obligation, that although gratitude leaves us latitude with respect to how we reciprocate a gift or favor, it also makes a very specific set of demands with respect to our future treatment of the benefactor: all ways of wronging or acting against the interests of a benefactor are ungrateful.[17] This point is useful to a theory of political obligation because it makes clear how gratitude can make definite *and* unlimited demands, and these are precisely the demands gratitude must be capable of making to explain how one could have a reason of gratitude to obey every law.

But this point is potentially devastating to a theory of betrayals of trust in terms of gratitude. If gratitude for trust behaved in the way that Walker claims gratitude for gifts or favors behaves, then *any* wrong to someone who trusts you would be ungrateful for that trust, not just the wrong they trust you not to commit. Yet only the wrong they trust you not to commit

17. A.D.M. Walker, "Political Obligation and the Argument from Gratitude," *Philosophy & Public Affairs* 17 (1988): 191–211.

counts as a betrayal of that trust. The challenge for my account of betrayals of trust as ingratitude for trust, then, is to explain how gratitude for trust can demand something definite—namely, that you do what you are trusted to do—without demanding anything else.

In reply, I claim that what gratitude demands can vary with the nature of the benefit. We can determine what gratitude for a given kind of benefit demands by asking what appreciation for that benefit necessarily motivates or inspires us to do. I think that Walker is correct that the demands of gratitude for gifts and favors are indefinite when it comes to making an appropriate return. Appreciating a gift or favor does prompt us to reciprocate. But this prompting is indefinite: there is nothing specific that appreciation for the gift or favor necessarily prompts us to do. Although this appreciation leaves us latitude with respect to reciprocation, it does move us not to wrong the benefactor or act contrary to his interests. The demands of gratitude for gifts and favors, then, are indefinite with respect to returns, but both definite and unlimited with respect to our future treatment of the benefactor, as Walker claims.

Now if the demands of gratitude for trust had these same contours, the account of betrayals of trust in terms of gratitude would fail. But trust is a different kind of benefit from a gift or favor. It is an honor. Benevolent motives underlie gifts and favors; an honor is just an honor. (A compliment may express an honor, but gratitude for a compliment is like gratitude for a gift or favor: it is a response to the benevolent motive underlying it. Trust is not a compliment, just an honor.) When we ask what gratitude for an honor demands by asking what appreciation for an honor necessarily inspires us to do, we discover that gratitude for an honor issues very different demands than gratitude for gifts and favors.

Whereas appreciation for a gift or favor prompts us to reciprocate, albeit indefinitely, appreciation for an honor does not seem to prompt us to make anything that could properly be called a "return" at all. Yet there is something very specific that our appreciation of an honor prompts us to do, namely, to live up to or prove ourselves worthy of the honor. Suppose that someone gives you a job. This could be a favor, if you were picked because you need the job, or it could be an honor, if you were picked because you were thought to be the one who would do the job best. If it were a favor, trying to do the job well might be the most appropriate return, but not necessarily. Perhaps you notice that the employer does not care all that much about how well the job is done, but

could really use a racquetball partner in the mornings. In that case, offering one's services as a racquetball partner might be a perfectly suitable way of discharging the debt. But if the job were given as an honor, appreciation of that honor would necessarily motivate you to do the job well. If you were to not take this honor as giving you extra reason to do good work, then you would have not appreciated the honor, even if—perhaps especially if—you instead were to take it as a reason to offer your services as a racquetball partner. Gratitude for an honor therefore appears to demand something specific, namely, that we strive to be worthy of the honor. When we trust, we honor the trustee in a limited way: we judge the trusted person as good willed toward us with respect to a particular obligation. Someone who appreciates the honor in trust will accordingly be moved to live up to this honor by displaying good will.

Gratitude for an honor, then, makes a specific demand. But does it only make one demand? Unlike appreciation for a gift or favor, appreciation of an honor inspires us *only* to live up to that honor—it has no other implications for our future treatment of the person who confers the honor. Suppose that after being hired as an honor, you did the job as well as you could, but felt in no way inclined to do your employer favors outside of work, even small ones. The employer might regard you as selfish and unhelpful, but could not cite the honor she had bestowed in hiring you as having given you a reason to perform the favor, as she might if she had hired you as a favor rather than an honor.

B. Objection: Entitlement to Trust

In some cases, it is natural to say that someone has *earned* our trust in virtue of their track record with us, or that someone is otherwise *entitled* to our trust because they have a solid reputation or because there is no apparent reason to doubt them. Perhaps someone has always proved trustworthy in the past about keeping my secrets and has, in this sense, earned my trust. I trust him to keep another secret, but on this occasion, inexplicably, he succumbs to the temptation to gossip and reveals it to others. Surely this is still a betrayal of my trust. If so, then according to the account of betrayals just given, spilling the secret should count as an ungrateful refusal to live up to the honor in my trust. But how can one be ungrateful for something that one has earned or is otherwise entitled to? This would imply that we should sometimes be grateful for what we are

owed, but it is not only not required but also demeaning to regard ourselves as placed in the debt of someone who is merely giving us our due.

I concede that trust can be owed, but I do not think this rules out the appropriateness of gratitude. The reason is that sometimes we should be grateful for what we are owed if we do not yet *deserve* it.[18] What someone is owed or entitled to, in the way of trust, depends on the evidence. But what someone deserves, in the way of trust, is settled by facts: since trust is prospective, the relevant facts are about their future performance. In other words, whether someone is entitled to our trust depends on the reasons available to us supporting his trustworthiness; whether someone deserves our trust depends on what he does *after* it is given.[19] So even if the reasons favoring someone's trustworthiness are overwhelming, while this may entitle them to trust, they still cannot be deserving of trust when it is given. The fact that the issue of desert is prospective in this way makes it appropriate to feel grateful for trust even when one is, in virtue of the evidence, entitled to it. This is not to deny that, in many cases, we need not or even should not feel grateful for what we are owed. It might be demeaning to be grateful for a paycheck. But the paycheck is also, unlike trust, already deserved.

C. Objection: Unwelcome Trust

Another potential challenge points out that trust can be experienced as a burden. If this were so, then there would appear to be nothing to be grateful for and no reason of gratitude to live up to this unwelcome trust. I do think that trust can sometimes be experienced in this way, but unwelcome trust poses no threat to my account. To show this, I will need to explain when trust can and cannot be betrayed.

Although all trust of the kind that has been my focus is marked by a readiness to feel betrayed in reaction to the trustee's failure, this feeling is not always *appropriate*. Only when trust is appropriate is it vulnerable to an actual betrayal of trust. One condition on appropriate trust is that the trustee must be under the obligation that the truster regards him or

18. An athlete may be entitled to a large signing bonus on the basis of his projected performance. But he does not yet deserve it. Gratitude for that bonus seems both intelligible and consistent with self-respect in a way that gratitude for compensation does not.

19. See David Schmidtz, "How to Deserve," *Political Theory* 30 (2002): 774–99, on the general idea of deserving X on the basis of what we do *after* receiving X rather than before.

her as being under. If you confusedly think that I have an obligation to you that in fact I lack, then you might trust me to carry it out and be ready to feel betrayed if I do not. But my failure to do so is not a betrayal of trust. It is only (incorrectly) liable to be experienced as such.

Another important condition on appropriate trust is that the trustee needs to have had a reasonable opportunity to reject the trust and to have declined it.[20] Even if I am already under an obligation to you, I may not welcome your trust. If I say, "Don't trust me," and thus reject your trust, you may stop trusting me. But there is no guarantee you will stop. You might not take me seriously, or perhaps you could not help yourself. The only effect my rejection is guaranteed to have (at least if I did not previously invite or accept your trust) is to make your trust inappropriate and remove the possibility of betrayal.

The fact that rejected trust does not generate a reason to perform as trusted is precisely what we would expect if the reason were a reason of gratitude. If you reject a benefit and it is forced on you anyway, it creates no debt of gratitude. Since on this account trust that has been rejected is inappropriate, in the sense that it is not vulnerable to betrayal, an explanation of betrayals in terms of gratitude remains viable.

But sometimes when a person finds trust unwelcome and has a reasonable opportunity to reject it, they decline to reject it. They may not want to alert the truster to the fact that the truster's expectations about them are mistaken. Asking your spouse to cease trusting you to be faithful because you find being trusted oppressive may eliminate the debt of gratitude (though not the original obligation). But it may also make it much harder to cheat, if that is your plan. On my account, since the trust has not been rejected, it remains appropriate and a betrayal is possible. Intuitively this seems correct. Yet the trust is unwelcome, and (so the objection goes) something that is unwelcome cannot create a reason of gratitude.

However, when we think about ordinary gifts or favors, we can see that the fact that a gift or favor is unwelcome does not necessarily rule out a debt of gratitude. If you offer to take me to the opera and I accept from politeness, I incur a debt of gratitude to you even though I secretly

20. This requires at least that the trustee knows or should know that she is trusted. If the trustee could not have been expected to know this, then even though someone may unreasonably persist in feeling betrayed by the trustee, there is no genuine betrayal.

despise the opera and wish you had not invited me. The reason is that I had a chance to reject your generously intended offer, and I declined to take it. Similarly, even when trust is unwelcome, if it is accepted, it generates a debt of gratitude that can be discharged only through doing what one is trusted to do. So again, although there are cases where unwelcome trust can clearly be betrayed, this is not a problem for an account of betrayals in terms of ingratitude because it is in precisely those cases that it is plausible that the unwelcome trust engenders a debt of gratitude.

VI. LYING AND BETRAYALS OF TRUST

Recall Solomon's suggestion that the lie is special because it commits a betrayal of trust. There was little point in pursuing this suggestion without an account of betrayals of trust, but now we have rectified this gap. The question now is whether this account of betrayals of trust can be used to vindicate Solomon's claim. Clearly a lie can betray trust. Suppose that Sue credits Bob with honesty, and thus expects him not to lie to her to get out of trouble, and Bob, knowing this, still lies to her about the coffee spilling. Sue's trust honors Bob, and he ungratefully repudiates that honor by failing to live up to it. His lie betrays her trust. But the fact that a lie *can* betray trust obviously falls short of establishing that the lie is special as compared with covert methods of deception.

It is clearly possible for a deceiver to use a covert method to deceive someone who does not trust him. (This is precisely why deceivers without credibility resort to such methods.) When the victim does not trust the deceiver not to commit a wrong, the wrong cannot betray trust. But it also seems possible for a deceiver to use a lie to deceive someone who does not trust him. Sue might expect Bob not to lie about what happened to the computer only because she thinks he is too unimaginative to make something up or too afraid of getting caught. She will still believe his lie, but he cannot be accused of betraying her trust, since she does not trust him not to lie, but only expects him not to lie for reasons that give him no moral credit.

Perhaps this overlooks the fact that the lie "invites trust," the distinguishing feature of the lie. I have not yet said what this invitation to trust consists in beyond characterizing it as an invitation to base a belief on the speaker's word. But let us suppose for the moment that the lie invites

the kind of trust that is vulnerable to betrayal. How might this be relevant? One possibility is that an invitation to trust in this sense inspires trust in those who lack it, making a wrong that would not have otherwise betrayed trust, for lack of trust, a betrayal. But it is hard to see why making such an invitation would inspire trust in someone who did not already see the speaker as trustworthy. If Sue does not already regard Bob as too honest to lie, why would Bob's inviting her to form a belief on the basis of his honesty lead her to do so? An invitation to believe a speaker on the basis of his honesty is not evidence for his honesty.

There are other possibilities, but they need not detain us: it is a mistake to think that the lie invites or, as Solomon says, "necessarily presupposes" the kind of trust that is vulnerable to betrayal. The lie invites the audience to base a belief on the speaker's word, and one can take another's word for something, thus fully accepting the invitation extended by the lie, by basing one's belief on the speaker's sincerity and knowledge. An expectation of sincerity is just an expectation that the speaker will say what he believes; one can expect sincerity from a speaker one does not regard as honest when the speaker has a nonmoral incentive to tell the truth. Thus, one can accept the speaker's invitation to form a "trusting" belief without trusting the speaker in the way that is vulnerable to betrayal. (I will continue to describe the lie as inviting trust, but it may be kept in mind going forward that the trust it strictly invites is not the kind that can be betrayed.)

So not all lies betray trust. Still, if a class of contexts can be identified in which lies would betray trust and covert methods would not, where the difference is explained by the fact that lies invite trust, then the lie would still count as special. But it turns out that not even this more modest goal is achievable. Here is the problem. Covert methods may not invite trust in any sense whatsoever, but trust need not be invited for a betrayal of trust to occur. All that a betrayal of trust requires is that there be an obligation that one is trusted to keep, and that the trust not be rejected. To illustrate, Sue may trust Bob not to deceive her, via any method of deception, to get out of trouble. That is, she may expect him not to deceive her about such things because she takes him to be trustworthy with respect to that general obligation.[21] If Bob knows this and

21. Sue's trust may be more selective than this, either because she thinks lying is harder to justify than noncommunicative methods, or because she thinks that Bob believes this.

has not rejected Sue's trust, then his staining of the cat's paws betrays Sue's trust even though this method of deception does not invite her trust. If an invitation to trust does not always set up a betrayal of trust, and if the absence of an invitation to trust never forecloses a betrayal of trust, it is hard to see how an account of the specialness of the lie in terms of betrayals of trust could possibly succeed.

VII. ABUSES OF TRUST

From my account of betrayals of trust as a form of ingratitude, no *method* of deception stands out as special. The lie can betray trust. But this is not because its invitation to trust sets up the betrayal; it is because a lie can count as a failure to do as one is (appropriately) trusted to do. And the relation to trust that a betrayal consists in—failing to perform as trusted—is a relation that any method of deception, regardless of its distinguishing features, is equally capable of bearing to trust.

However, a failure to perform as trusted is not the only kind of morally problematic relation a method of deception can bear to trust. Some forms of deception may also make an improper *use* of trust. An improper use of trust, I will argue later, differentiates the lie from covert methods of deception. The invitation to trust may not set up a betrayal, but it does mark an intention to use trust to deceive someone. But this problematic relation, which I will call an "abuse of trust," stands in need of an account just as much as the relation of a betrayal of trust.

A. Use of Commendable and Charitable Attributes

It will be useful to begin with some general thoughts about the moral significance of "use." I propose to take the notion of "using a person" very broadly: as taking an opportunity that is created or sustained by one

---

An interesting implication is that beliefs about the specialness of lying relative to other methods can, to some extent, be self-fulfilling. The kind of trust that is vulnerable to betrayal requires regarding the trustee as under an obligation. If Sue thinks that covert methods of deception are permissible to get out of minor scrapes, but not lies, then even if she is mistaken about this, it follows that Bob cannot betray her trust by staining the cat's paws, but he can betray her trust by lying to her. But it is not the fact that the lie invites trust that explains this difference, but rather Sue's beliefs about the relative permissibility of various methods of deception.

or more of the person's attributes. Virtually every way of treating a person takes an opportunity that rests on one attribute or another in the person, so it should be clear that the notion of "using a person" I have just defined is not meant to capture our pejorative notions of "using a person" or "exploiting a person."

Although the mere fact that a form of treatment of a person also "uses the person" in this sense seems unimportant, the identity of the attribute that was used can matter. There appear to be attributes that place the opportunities they support morally off limits, at least for certain purposes and by certain people.[22] To give one example, perhaps all of the opportunities that unrequited love presents are morally off limits to the object of that love. The cases that will concern me here, however, are those in which a form of treatment is wrong independently of the fact that a particular attribute was used, but in which the use of that attribute in particular appears to aggravate the wrong.

What kinds of attributes place constraints on how they may be used and by whom? There is a wide variety of attributes, from vices and character weaknesses to conditions like drunkenness and blindness, that lay us open to wrongs.[23] Even traits we regard as strengths create opportunities for wrongdoers: for example, a prudent person is a better target for a mugging than a rash person. I will not run through all the possibilities here, but for most attributes, the fact that a wrong deliberately took an opportunity opened up by the attribute in no way aggravates the wrong.

Matters are different with regard to two kinds of attributes: morally admirable attributes and charitable attributes. Morally admirable attributes include emotions and motives like compassion, benevolence, and a sense of responsibility. (Often these motives issue from an underlying virtue, but even when out of character, they still seem to place constraints on their usage.) Consider the difference between a scam that employs an investment proposal and a scam that employs a charity proposal. The investment scam may be playing on greed or ordinary prudence, but in either case, the fact that the success of the theft rests on

22. For the idea of a constraint, see the discussion in Alisa Carse and Margaret Olivia Little, "Exploitation and the Enterprise of Medical Research," in *Exploitation and Developing Countries*, ed. Jennifer Hawkins and Ezekiel Emanuel (Princeton, N.J.: Princeton University Press, 2008), pp. 206–45.

23. For a useful taxonomy, see Joel Feinberg, *Harmless Wrongdoing* (Oxford: Oxford University Press, 1988), pp. 176–210.

selfish or self-interested motives adds nothing to the wrong. But the charity proposal triggers the victim's sympathy or compassion for others, and the use of these motives seems to worsen the violation. This is not just because we feel that a compassionate person is less deserving of the wrong. The person who falls for the investment proposal may be a sympathetic person who would have bitten on the charity proposal had that been presented instead. Yet the deployment of a charity proposal as bait still seems worse than the deployment of the sham investment proposal.[24]

A wrong that takes an opportunity presented by a morally commendable attribute appears to display not only whatever disrespect for the person is expressed by the wrong as such, but also a distinct form of disrespect for the attribute. Such attributes are not to be misused. I will not say anything further about the nature of the reason underlying this constraint, except to note that it is a reason that applies to every potential wrongdoer, and that it applies only insofar as the attribute is morally admirable. (Playing on misplaced compassion or a distorted sense of responsibility does not aggravate the wrong.) These two features distinguish the constraints associated with commendable attributes from those associated with the next category of attribute, charitable attributes.

Charitable motives include compassion and generosity, so it is quite possible for an attribute to be both commendable and charitable. But what makes an attribute "charitable" in the relevant sense is just that it is *favorable* toward someone. A willingness to help someone out shows favor in this sense, even when helping that person is morally inappropriate in the circumstances. Furthermore, the constraints it imposes on its usage apply only to the person who is the object of the favorable attitude, not to any potential wrongdoer.

To illustrate, return to the example of the scam that plays on compassion. As mentioned, there is a difference between a scam that plays on

24. No doubt part of the reason we condemn the charity scam more harshly is on account of its broader consequences. Abusing someone's compassion tends to discourage future displays and if nothing else squanders this particular display. Thus, not only the mark but also those who might have otherwise benefited from his generosity are victimized. But this cannot be the whole explanation. Even if there were no chance of discovery, and even if the victim stood ready to purchase a yacht had his generous impulses not been stirred by this particular scam, the sense that the wrong is graver persists.

compassion and one that plays on a self-interested motive, whether derived from greed or ordinary prudence. But there is a further difference between a scam that plays on compassion toward others and a scam that plays on compassion toward the wrongdoer himself. The victim of the latter scam would rightly experience the fact that the wrong preyed on her compassion *for* the wrongdoer as an added insult. And this is so quite apart from whether the compassion is appropriate or admirable in the circumstances. Perhaps the victim might be accused of being generous with the scammer to a fault, given the victim's other responsibilities. But this only seems to enhance the wrong done to the victim, rather than to detract from it.

Although the nature of the reason behind the constraints imposed by commendable attributes may be difficult to pin down, things are otherwise in the case of the constraints imposed because an attribute is charitable. The reason is gratitude. What other kind of reason would apply only to the object of the attitude and because the attitude is charitable toward that person? Whereas one should be respectful of the opportunities presented by morally admirable attributes, one should be *grateful* for the opportunities presented by charitable attributes. The demand that these opportunities make is not that they be reciprocated by returning good for good, or even by avoiding returning bad for good. Instead, what they demand is that they not be *used* to return bad for good. (This is why a scam that plays on someone's self-interest is not as odious as one that plays on her compassion. The opportunity to play on her compassion was there, but the scammer declined to take it.) Using such an opportunity to wrong the person is an ungrateful misuse of that attribute.

B. Use of Trust

Now the question is whether the use of trust can also, like the use of the motives I have mentioned, make a moral difference. Although trust is an expectation, not a motive, it does afford a variety of opportunities to others, especially to the one who is trusted. Trust involves an expectation that the trustee will perform as trusted, and this expectation provides the trustee with opportunities for contact with the truster and with what the truster cares about that the truster might otherwise regard as too risky or too costly. The truster may offer or agree to hire the trustee, let her do her

work without supervision, let her borrow her car or her money, crash at her apartment, come to dinner, go on a date, or look after her children.[25]

Among the opportunities trust characteristically makes available to the trustee are opportunities to wrong the truster or to wrong the truster more conveniently, safely, or gainfully. It is a very familiar fact about trusting someone that it makes one vulnerable in these ways. Trust prompts us to create opportunities to wrong us that were not already there, such as when we confide in someone or lend them money, and inhibits us from closing or burdening opportunities to wrong us that were already there, such as when we decline to check up on someone's work or secure our valuables.

Is trust an attribute that imposes constraints on its usage, and if so, is this because it is morally admirable or because it is charitable? (Or both?) I do not want to deny that trust can be morally admirable,[26] but I think this is the exception rather than the rule. If it were morally admirable as a rule, then since its constraints would apply to all wrongdoers, we would expect to find wrongs that take opportunities based on trust in someone other than the wrongdoer as special. Such wrongs are not infrequent, especially on the Internet, where scammers impersonate someone the victim trusts or pretend to occupy a trusted role to get the victim to release information. But it does not seem to me that these wrongs have the same character as misuses of a victim's compassion or benevolence for someone else.

Rather, given what I have already said about trust, we should expect any constraints it imposes on its usage to derive from the fact that it is a charitable attribute. Although the attitude of trust is not, like generosity, aimed at the interests of its recipient, it does honor its recipient, and as

25. When someone is trusted, the expectation that they will not commit a wrong is supported by the thought of their trustworthiness. Similar expectations about a person's behavior can be formed on other bases, for example, the thought that they lack the opportunity to commit the wrong or have nothing to gain from committing the wrong or have strong incentives not to commit the wrong, such as a fear of exposure and sanctions. But such expectations arise only in a narrower set of circumstances: either where the truster is already assured that there is a lack of opportunity or incentives against the wrong, or where the trustee is willing and able to take sometimes costly precautions to give themselves that assurance.

26. The cases where trust seems morally admirable tend to be those where there are reasons for doubt, but the expectation of trustworthiness is nonetheless formed or sustained by the exercise of some morally admirable motive or emotion. In such cases, the trust is praiseworthy because it displays a praiseworthy motivational attribute.

I have already argued, this feature can ground reasons of gratitude. These reasons will apply only to the recipient of that trust: the attitude is charitable only toward the recipient after all.

So does a wrong that abuses trust seem ungrateful, just as a scam that plays on someone's compassion for oneself seems ungrateful? I believe that it does, but there is a difficulty in isolating the source of that ingratitude. There is something ungrateful about my taking an opportunity presented by my spouse's trustful lack of monitoring to conduct an illicit affair. But since I am also committing a betrayal of trust by failing to perform as trusted, there is already an element of ingratitude in the wrong that is in place independently of the origin of the opportunity I took. Does my *use* of her trust to facilitate the betrayal add anything?

It does: it compounds the ingratitude. To see this, it helps to notice that there are abuses of trust that are not betrayals, and betrayals that are not abuses. These forms of ingratitude often co-occur, but they can come apart. Suppose that you believe your neighbor is immune to temptations to steal, yet is prone to snooping and telling secrets. Suppose too that your neighbor is aware of your beliefs about her. You might invite this neighbor to house-sit for you while you are on vacation, having taken care to conceal any evidence of embarrassing personal habits, while leaving your valuables out in the open and unsecured. If you inadvertently forget to put away some embarrassing evidence, and if your neighbor spots it and spreads the news, then your neighbor will have wronged you. However, you cannot complain that she has betrayed your trust. She has not stolen anything, and that was all you trusted her not to do. But it seems to me you can complain that your neighbor has ungratefully abused your trust. She is in your house because you trusted her around your valuables, and someone grateful to be trusted around your valuables should be reluctant to use any of the opportunities flowing from that trust against you, even those that were unwittingly provided.

Betrayals of trust without abuses are also possible. Suppose that you trust me not to disclose certain kinds of personal information about you to strangers, and hence would have given me this information if I had asked, but instead I decide to find it out the hard way, by rifling through your garbage. Certainly I betrayed your trust, but I avoided abusing it because I refrained from taking an opportunity provided by your trust. My motives for eschewing opportunities supported by your trust may not always do me credit. That is, my reason for acquiring the information

the hard way might be to keep you from knowing what I did and might have nothing to do with a desire to avoid compounding the ingratitude of the wrong I commit. Nevertheless, by avoiding the use of such an opportunity, I do avoid displaying an additional form of ingratitude.

VIII. DECEPTION AND ABUSES OF TRUST

Successful deception, like any way of having an effect on a person, takes an opportunity that rests on one or more attributes of the person. Just as earlier it was pertinent to ask what motivational attributes may be exploited by a scam, now it is pertinent to ask what epistemic attributes may be exploited by the deception that is part of a successful scam. Once again, we are confronted with a wide variety of potential attributes: traits, like gullibility and susceptibility to wishful thinking, as well as curiosity and open-mindedness; general conditions, like senility and naïveté; gaps in one's background knowledge, large or small; and so on. But none of these, in my view, places constraints on their usage.[27] Nor should they be expected to, since none is morally admirable or charitable. Trust of course is different. Just as using trust to commit a theft aggravates the theft, so we should expect a deception that uses trust to aggravate the deception. But when does a deception that betrays trust also deliberately take an opportunity based on trust? In other words, when does a deception *abuse* someone's trust?

A. Methods of Deception: Lying

Recall Bob and Sue. Let us assume that Sue trusts Bob not to deceive her via any method, so that he commits a betrayal of trust with whatever method he employs. But there is now the additional category of an abuse of trust to consider. Even when different methods would betray trust, one method could still be special because it also abuses that trust, violating an additional demand of gratitude. To address this question, I will need to be rather more precise about the differences between various

27. Deceit that takes an opportunity made available by an attribute that is a misfortune for the person, like senility, may seem worse than other deceit. But this may be because the person is a less-eligible target for a wrong, in virtue of their disadvantage. In other words, it may be that what makes the wrong worse is that it is a form of piling on, not that the person's misfortune was *used*.

methods of deception than I have been up to this point. In Grice's famous article "Meaning," he distinguishes three ways of getting someone to believe something.[28] I adopt his analyses in what follows, with a few refinements.

One way of getting someone to believe something is by "telling." When Bob tries to get Sue to believe *p* by telling her that *p*, he not only intends her to recognize that he intends her to believe *p*, but also intends her recognition of his intention to be the *basis* for her belief. This is the sense in which he invites her to take his word for *p* or invites her trust. (More specifically, her recognition of his intention is to be the basis of her belief because she takes it to be an intention to produce a belief in her that he *shares*, that is, is expressing sincerely.) Grice labels another way of inducing a belief "merely getting someone to think." When Bob merely tries to get Sue to think that *p*, not only does he not intend her recognition of his intention to serve as the basis for her belief, but he does not intend her to recognize his intention to get her to believe *p* at all. Finally, there is a way of getting someone to believe that Grice calls "deliberately and openly letting someone know," or less cumbersomely, "showing." When Bob shows Sue that *p*, he does intend her to recognize his intention, but he intends what he shows her to be the basis for her belief, not the recognition of his intention.

As should be apparent from the earlier discussion, the method of deceit that involves telling is the lie, and the method corresponding to merely getting to think is what I have been calling "covert deceit." My first task, now that we are equipped with analyses of these methods, is to determine whether appeal to the notion of an abuse of trust can distinguish lying from covert deception. Following that discussion, I will introduce a third category of deceit, corresponding to showing, in order to compare it to lying. This will help to determine whether lying is truly special, from the point of view of trust, or just an instance of a more general method that is special.

When Bob lies to Sue about the cat to get her to believe that the cat spilled the coffee, he intends his utterance to make her aware of his intention that she believe this, and further, that her recognition of this intention be the basis for her belief. How exactly is this supposed to work? There is more than one way in which recognizing someone's

28. Paul Grice, "Meaning," *Philosophical Review* 66 (1957): 377–88.

intention to produce a certain belief can provide a reason to form that belief.[29] But the way the belief will come about in someone who trusts Bob, as Sue does, is as follows. First, she will take his intention to get her to believe *p* as a reason to believe only if she interprets it as an intention to produce a belief he shares: that is, an intention to *inform* her. Secondly, what will prompt that interpretation in Sue is her trust, the thought that he would not try to get her to believe something he disbelieves or does not believe because he is honest. This is the way in which Bob's plan, in lying to Sue, involves the activation of her trust: it is her expectation of honesty that will lead her to interpret his intention as an intention to inform, and it is only if his intention is given this interpretation that it can serve as a basis for her believing *p*.

That Bob deliberately took an opportunity resting on Sue's trust does not imply that her trust was necessary for the success of his deception. Suppose that, even apart from Bob's honesty, Sue has backup reasons to think that Bob would only intend her to believe what he himself believes. Perhaps she thinks the consequences of getting caught in a lie are so much worse than the consequences of having ruined the laptop that even if he were not honest, he would still not try to deceive her. What this means is that Bob would still have had an opportunity to deceive her with his lie even if Sue did not trust him. But it remains the case that the *actual* opportunity he takes rests on her trust, and this is what it is to abuse her trust. Thus, lies that betray trust will always abuse that trust, violating two distinct demands of gratitude.

B. Methods of Deception: Covert Deceit

Next consider Bob's framing of the cat, an example of covert deceit. As with the lie, Bob is introducing a change in Sue's environment with the intention of leading her to form a certain belief. But there are two differences. One difference is that he wants her to treat the cat's paws as *direct evidence* for the cat's guilt. In other words, the cat's paws are supposed to be the basis for her belief, not her recognition that he intends her to believe that the cat did it. The second difference is that not only does Bob

29. If Sue thought Bob was reliably mistaken about certain things and also intended to deceive her about these things, then her recognition of his intention that she believe *p* would give her a reason to believe *p*. But this is not the way Bob would expect the recognition of his intention to work in someone who trusts him, and Sue trusts him.

not intend Sue to believe on the basis of recognizing his intention that she believe something, but he does not even want her to be aware that he possesses such an intention. This is not just a matter of his not wanting her to know that the cat's paws were a product of such an intention, which would obviously spoil its value as direct evidence. Rather, Bob does not want her to know he has any designs on her beliefs at all. (It is obvious why a deceiver who is mistrusted would not want to advertise his intentions, but even a deceiver who is trusted by his target, as Bob is, may want them concealed for the sake of plausible deniability should the victim subsequently discover the truth.)

Does Bob, in covertly deceiving Sue, deliberately take an opportunity resting on Sue's trust? When Bob lies, he is trying to activate Sue's expectation of his honesty. Whereas, if everything goes as planned with the cat's paws, the thought of him will not even cross her mind. She will just notice the cat's paws on her own and the evidence will convince her. Insofar as Bob's plan is to get Sue to believe without even thinking of him, much less his honesty, he does not appear to be deliberately taking an opportunity resting on her trust. Where the deceiver's behavior is supposed to be the evidence, such as in Kant's example of someone's packing their bags to pretend they are going on a trip,[30] the thought of the deceiver obviously needs to cross the victim's mind. But again, the target's trust will be activated only if the thought occurs to her that he may be trying to get her to believe something, and this is not a thought he is trying to produce in her. In fact, it is quite the opposite.

So far it looks like covert deceit does not abuse trust even though it may betray it. But there is a potential complication. Despite Bob's best efforts to avoid prompting any thought of himself in Sue, he may fail; but he is nevertheless confident of success in deceiving her because the thought of him will activate her trust. Perhaps Sue had been recently tricked and is very suspicious of her environment, or perhaps something else makes her mindful of the fact that Bob could have arranged this scene and had an incentive to do so. If such a thought does occur to Sue, then she may continue to take the cat's paws as direct evidence only because this thought also activates her expectation of his honesty and leads her to dismiss his possible involvement. If Bob knows there is a

30. Immanuel Kant, *Lectures on Ethics*, trans. Peter Heath, ed. Peter Heath and J. B. Schneewind (Cambridge: Cambridge University Press, 1997), p. 426.

good chance she will entertain the thought of his involvement, he may even regard the route that activates her trust as the route his deception is most likely to take.[31]

If Bob's deception does succeed via the activation of Sue's trust, then his deception takes an opportunity supported by her trust. However, he is still not *deliberately* taking this opportunity, as he does with the lie. Bob may know that if his efforts to slip the deception past her without his crossing her mind fail, he will need her trust to be activated if his deception is still to succeed. He may even think that in all likelihood this is how his deception will succeed. But this does not mean that he is trying to activate her trust; indeed, he is still trying to avoid activating it.

Thus, when covert deception betrays trust, it does not also abuse that trust. Since lies that betray trust also abuse it, lying appears to be special as compared with covert deceit, at least in contexts where both methods would betray trust. This solves the puzzle I presented at the outset of the article, which was to explain how the distinguishing feature of the lie, which makes it less manipulative than covert deceit, can somehow make it worse.

C. Methods of Deception: Deceit by Showing

There is a further question I wish to address. The lie invites trust, which implies that the intention to produce a belief is meant to be recognized, and that the recognition of this intention is meant to be the basis of the belief. Are both of these features needed to explain why it abuses trust, or not? If not, then even though the lie is special as compared with covert deceit, it may not be the only method that is special in that way. I will explore this possibility by thinking about the third way of getting someone to believe that Grice identifies: *showing*.

Grice illustrates *deliberately and openly letting someone know*, or *showing* as I am labeling it, with the example of Herod's presenting John the Baptist's head on a platter to induce Salome to believe John is dead.

31. It is also conceivable that even if she realizes that he stained the cat's paws, her trust might still enable the deception to succeed by leading her to assume he must have stained the cat's paws to make the truth about what happened more evident to her. However, if she is also aware that he did not intend for her to know he stained the cat's paws, it is far more likely that she will conclude that he is trying to deceive her: trust can only be stretched so far before it breaks.

To make this a case of deception, suppose that it is John's twin's head on the plate. As with covert deceit, the victim is intended to believe on the basis of the direct evidence. It is the severed head that is supposed to do the convincing. But unlike covert deceit and like lying, the target is meant to recognize that the deceiver is trying to produce a belief, usually by means of some gesture accompanying the evidence.

Sometimes the evidence will make the belief completely obvious. The intention to produce a belief may be overt, but its recognition is entirely unnecessary for producing the belief: the evidence can do all the work on its own. If Salome knows what John looks like and the question of whether he is dead is on her mind, just seeing John's twin's head will produce the false belief, whether or not she also recognizes that Herod is showing it to her. I will call this "deceit *with* showing" to indicate that the act of showing accompanies but does not help accomplish the deceit. Deceit *with* showing is different from covert deceit, since the intention to produce a belief is overt. But it employs exactly the same *means* of deceiving the target as covert deceit, since the evidence makes the belief completely obvious. Deceit *with* showing may, unlike covert deceit, activate trust since the victim is meant to recognize the intention to produce a belief. But since the evidence makes the belief completely obvious, one cannot be accused of deliberately taking an opportunity supported by the activation of trust. Neither the recognition of the intention nor the trust it triggers plays any role in bringing about the belief.[32]

But if the evidence, instead of the recognition of the intention, is supposed to do the convincing, how could showing ever help produce a belief? One might wonder whether deceit *by* showing, as opposed to deceit *with* showing, ever occurs. But it does, and the reason is that the recognition of the intention can help produce the belief without in any way serving as the basis of the belief. Sometimes the recognition of the intention may be needed to get someone to *notice* the direct evidence. Sue might not notice the cat's paws on her own. But if Bob points to the

32. Now one's reason for making one's intention overt even though the evidence makes *p* completely obvious may be that one wants to get them to believe that one also believes *p*. If so, then the deceit may still abuse trust, since the recognition of one's intention to get them to believe *p* or to believe that one believes *p* will be effective only if it is taken as an intention to inform, and it is their trust that will support this interpretation. But in such a case, one is not deceiving them that *p* by showing; rather, one is deceiving them into believing that one believes *p* by showing.

cat, making her aware that he wants her to come to believe something by looking at the cat, she may take herself to have a reason to examine the cat carefully and may thereby come to notice its stained paws. The cat's paws, and not Bob's intention, are still intended to be the basis for her belief, but the recognition of his intention may be needed to prompt her to notice them. And even if Sue independently notices the stained paws, some *processing* on her part may still be necessary to reach the desired conclusion, and processing demands effort. Again, knowing that Bob has a belief in mind for her may be necessary to motivate her to process her way to that belief.[33]

So even when a deceiver is counting on the evidence to do the convincing, not the recognition of his intention, he may still need the target to recognize his intention in order to *incentivize* her to notice and process that evidence. And when the deceiver intends the recognition of his intention as a means to producing a belief in these ways, rather than as a basis, he still abuses the target's trust. Just as Sue will take Bob's intention to get her to believe something as a basis for belief only if she takes it as an intention to inform her, so she will take this intention as an incentive for noticing and processing some direct evidence only if she takes it as an intention to inform her. If she were to take it as an intention to deceive her or to produce a belief he does not share, why bother noticing or processing? And since she trusts Bob, she will interpret it as an intention to inform her because the recognition of his intention will activate her trust. So when Bob counts on her recognition of his intention to motivate her to notice and process the evidence, he is deliberately taking an opportunity that flows from her trust, even though it is ultimately the evidence and not her trust that is supposed to be the basis of her belief. Deceit *by* showing, therefore, is morally on a par with lying.

IX. CONCLUSION

The fact that deceit *by* showing is on a par with lying implies something interesting about what category of deceit is morally fundamental, as far

33. For discussion of the ways that recognition of an informative intention can help lead to belief, see Dan Sperber and Deirdre Wilson, *Relevance* (Malden, Mass.: Blackwell, 1995), pp. 46–50.

as trust is concerned. Before I explain what this is, a brief review is in order. Commentators have tended to focus on distinguishing lying from falsely implicating, seemingly taking for granted that both are special compared with other methods of deceit. But it is far from obvious why the lie should be regarded as worse in some way than covert deceit, which lacks its forthrightness and seems more manipulative. An appealing answer is that lying is open because it invites trust, something covert deceit does not do, and the invitation to trust sets up some special trust-based wrong that the lie commits.

Such an explanation can hardly be convincing without some account of the trust-based wrong. But the moral significance of trust has been largely unexplored. I have tried to make some progress on this topic by linking the significance of trust to gratitude and its various demands. Whatever general value this approach to betrayals of trust might have, it does not help to explain what sets the lie apart from covert deceit. Yes, a lie can betray trust, but covert deceit is equally eligible to commit that kind of wrong. The account of abuses of trust, on the other hand, can help to vindicate the specialness of the lie compared with covert deceit. It does appear to be true that the lie's invitation to trust makes it special, and this explains how the openness of the lie, which may appear only to make it less manipulative than covert deceit, also has a role in making it worse.

Now it might be tempting at this stage to think that the fundamental difference between methods, as far as abuses of trust are concerned, is between methods that invite trust and those that do not. But since another method that does not invite trust, deceit *by* showing, turns out to be morally on a par with the lie, there must be a broader category of method that is more fundamental. This is the category of deceit that involves an overt intention to produce a belief, where the recognition of this intention is meant as a *means* to the belief, whether as an epistemic basis (as in the lie) or as a practical incentive to notice and process some evidence (as in deceit *by* showing). Lying and deceit *by* showing are on one side of the line. Covert deceit and deceit *with* showing are on the other, since neither uses the recognition of the intention as a means to producing the belief. Covert deceit does not even intend the recognition of the intention; deceit *with* showing does, but this recognition is not needed to produce the belief, as the evidence makes it completely obvious.

I can now state the morally significant contrast in more familiar terms. Although Grice would disagree,[34] intending that one's intention to produce a belief be recognized counts as intending to *communicate* that belief. Thus, the morally fundamental difference between methods of deception, as far as trust is concerned, is between deceit *by* communication and all other methods. The lie and deceit *by* showing both communicate the belief and produce the belief by means of that communication. Covert deceit communicates nothing, and although deceit *with* showing communicates a belief, it does not produce the belief by means of that communication. The lie may invite trust, but it is with the use of communication to produce a belief that the morality of trust is most fundamentally concerned.

34. Grice denies that in "deliberately and openly letting someone know" something, one (nonnaturally) means it. Grice, "Meaning," p. 218. But see Sperber and Wilson, *Relevance*, pp. 53–54; and Stephen Neale, "Paul Grice and the Philosophy of Language," *Linguistics and Philosophy* 15 (1992): 547–49.